Philosophers' Imprint

VOLUME 25, NO. 3
JULY 2025

Doing Harm While Allowing It: Towards a New Framework for the Doctrine of Doing and Allowing

David Turon

University of Massachusetts Amherst

DOI: 10.3998/phimp.3370

1. Introduction

Generally speaking, it is seriously wrong to do harm to others. It is also often seriously wrong to allow harm to others. Some nonetheless hold that doing and allowing harm are morally inequivalent. They endorse the Doctrine of Doing and Allowing (DDA): the view that it is harder to justify doing harm than *merely* allowing harm, all else being equal.[1] For example, it seems wrong to deflect a lethal threat onto an innocent in order to save oneself, but permissible to allow a lethal threat to reach an innocent in order to save oneself. The DDA naturally accounts for this. But others deny that there is any morally significant difference, arguing that when all else is equalized, doing harm is no worse (nor harder to justify) than allowing harm.[2]

This way of putting the issue pits doing harm against allowing harm. But why should we do that? Does one do some harm only if one does not allow it, and does one allow harm only if one does not do it? Are doing and allowing mutually exclusive? I do not think so.[3] Carefully attending to the ways in which doing and allowing harm can both come together and come apart, as I shall argue, reveals a surprising amplificatory relationship between doing and allowing harm. Further, explaining this relationship may clue us in to some of the reasons why most ordinary harm-doing, which involves *both* doing and allowing harm, is harder to justify than merely allowing harm.

Here is a roadmap for the paper. In §2 I argue that, when it comes to doing and allowing harm, we should distinguish between (i) *merely* allowing harm, (ii) *merely* doing harm, and (iii) doing harm *while* allowing it. On its own, surprisingly, doing harm seems to carry little or no moral weight; in fact, it seems *easier* to justify *merely* doing harm than *merely* allowing it. But doing harm *while* allowing it, which is what

1. For a few examples of defenses of the DDA, see Woollard (2015), Foot (1967, 1984, 1985), and Quinn (1989). See also Liu (2012), who offers a consequentialist rationale for the moral significance of the doing/allowing distinction.
2. See, e.g., Bennett (1995) and Rachels (1975).
3. Many authors writing on the DDA (e.g., Woollard, 2015) have been clear to avoid this commitment, focusing on the contrast between doing harm and *merely* allowing harm, where the latter is explicitly defined as allowing harm *without* doing harm.

happens in most cases we intuitively classify as *doing harm*, seems harder to justify than merely allowing it, as the DDA proclaims.

In §3 I argue that, while there may be independent moral weight associated both with doing and allowing harm, the moral weight of one or the other (or both) is *amplified* when the two converge: there is a moral synergy between doing and allowing harm. In particular, I argue that it takes much more to justify doing harm to someone while allowing them to suffer that harm than it does to *merely* do an equivalent harm to one while *merely* allowing an equivalent harm to another. This result gives rise to a puzzle: why is it so much harder to justify doing some harm while allowing *that* harm than merely doing some harm plus merely allowing some distinct but equivalent harm?

In §4 I sketch two distinct but compatible ways to resolve the puzzle. First, in §4.1 I explore the possibility that introducing a threat of harm generates a special duty not to allow that harm, thereby making the duty not to allow harm more stringent. Second, in §4.2 I explore the possibility that the rights-infringements involved in doing harm constitute graver wrongs as they become more harmful for the victim; this view is in line with Judith Jarvis Thomson's (1990) Aggravation Principle.

I then conclude in §5 by suggesting how we could recast the DDA to better reflect the relationship between doing and allowing harm.

2. Doing *Versus* Allowing Harm?

Some have distinguished standard instances of harm-doing from cases of allowing oneself to do harm.[4] For example, suppose someone poisons a victim's drink, but later has the chance to warn the victim before she takes a sip. In refraining from warning the victim, the poisoner allows himself to have done harm to the victim. Thus, the poisoner both does and allows harm to the victim. But rather than taking instances in which an agent both does and allows a harm to be a special kind of case, we should recognize that *most*, though not all, instances of what is ordinarily regarded as *doing* harm are also instances where the harm-doer *allows* the harm in question. To see why, let us first elucidate what it is to allow harm.

The rough, intuitive idea is this: you allow someone to be harmed in the event that they suffer harm and you could have done something such that, if you had done it, they wouldn't have suffered that harm.

More precisely, an agent allows harm of magnitude m to a subject if and only if (i) that subject suffers some token harm H of magnitude m, and (ii) the agent had an option available to her such that, had she taken it, the subject would neither have suffered H nor have suffered any harms of total magnitude m (or greater) instead of H that he didn't actually suffer.[5]

Here is an example. Suppose you trip on a log and break your leg, and you suffer some associated harms. Bob was able to remove that log from the path, but he didn't. If Bob had removed the log, you wouldn't have suffered the harms associated with breaking your leg, nor would you have suffered any new harms of equal or greater magnitude that

4. Judith Thomson has pointed this out (1985: 1400). There has been more recent controversy about the moral significance of allowing oneself to do harm (see, e.g., Hanna, 2014; Unruh, 2021).

5. I am indebted to an anonymous referee for pushing me to be clearer about 'allowing harm' here. In particular, we should take care to acknowledge the distinction between failing to avert a *token* harm, failing to avert harm of a certain *kind*, and failing to avert a *total magnitude* of harm. My proposed definition falls into the third category, although it mentions particular harms in the *definiens*. This proposal will ultimately deliver equivalent results regardless of how modally fragile we take token harms to be; that is, regardless of how much a harm's features can vary and still count as *the same harm*.
Here is an example to illustrate why. Suppose you have two options: option (a), on which Ann will get a bruise, and option (b), on which she will get a bruise at a slightly different place on her body but that causes or constitutes harm of exactly the same magnitude as that which she suffers on option (a). All else is equal. Whether we classify the harm arising from the bruise when taking option (a) as *the same harm* as that arising when taking option (b), you won't have allowed her to suffer any magnitude of harm. First, more simply, if Ann counts as suffering the same harms on options (a) and (b), then she would have suffered all the same harms no matter what you did, and so you won't have allowed her to suffer any magnitude of harm. If she doesn't count as suffering the same harms on options (a) and (b), then on option (b) Ann will suffer a different harm of the same magnitude as the one she suffers on option (a). So again, by my definition, you would not count as allowing her to suffer any magnitude of harm, no matter which option you take.

you did not actually suffer. Thus, Bob allowed you to suffer harm of a certain magnitude: namely, the total magnitude of the harms associated with breaking your leg.

With the notion of allowing harm so clarified, we can now see that there is a simple reason why, when someone does harm of some magnitude, they also usually allow harm of that magnitude. Consider the following case:

> *Button*: Ann presses a button. Once this button is pressed, it sends a current to shock Bob a few seconds later. Ann is unable to do anything to stop the shock from reaching Bob once she presses the button.

It is clear that Ann *does* harm of a certain magnitude to Bob by pressing the button. But she also simultaneously *allows* Bob to suffer that magnitude of harm, for at the moment she pressed the button she also had the ability to do something else instead (e.g., sit still) such that, had she done it, Bob would not have suffered a harmful shock soon after. But in *Button*, Ann loses the option to avert the harm Bob suffers the instant she initiates the threat.

Here is the upshot: even if one cannot undo a threatening sequence at any moment after it is initiated, it can still be true that one allowed someone to suffer some magnitude of harm from the very harms produced by the threatening sequence. This is usually because one could have refrained from launching the threatening sequence in the first place. I submit that, for this reason, most instances of doing harm are also instances of allowing harm — not merely in virtue of the harm-doer's having an ability to cancel a threat *after* it is launched, but in virtue of the fact that, at most times when someone does harm to a victim, (i) she had the ability not to launch the threat of harm in the first place, and (ii) had she not launched the threat in the first place, the victim would not have suffered that harm (or harm of an equivalent magnitude instead).

Why, then, might someone be inclined to describe Ann as doing harm to Bob, *rather than* allowing harm to him, in cases like *Button*? We might think that this is because the meanings of ⌜doing ϕ⌝ and ⌜allowing ϕ⌝ are such that one entails the negation of the other. But if that were so, I could simply say that the morally relevant notions I refer to by 'doing harm' and 'allowing harm' are distinct from the referents of those expressions in ordinary English. Nonetheless, I think that 'doing' does not preclude 'allowing', nor vice versa, even in the ordinary sense. I suggest that describing someone as allowing harm merely conversationally implicates that she only allowed harm without doing it. Again, most of the time when we do harm, we also thereby allow it. Normally, then, it goes without saying that if one does harm, one also allows it — it typically goes without saying, in other words, that if you did harm to someone, you could have done something else such that, had you done it, the victim would not have suffered harm. Given this point, describing someone as allowing harm conversationally implies that she did not do harm, as omitting that further piece of information would run afoul of the Gricean Maxim of Quantity.[6] And this would explain why we might be inclined not to describe Ann as allowing harm in *Button*.

To support this diagnosis, observe that we can cancel the relevant implicature flowing from describing someone as allowing an outcome, using our ordinary sense of 'allow'.

> (1) Did you allow the cat to go outside?
>
> Yes. In fact, I put the cat outside myself.
>
> (2) Did you allow the bomb to go off?
>
> Yes. In fact, I detonated it myself.

To my ears these exchanges check out.

Note again that it would not be a problem if the notion of *allowing harm*, as I have defined it above, did not fully comport with ordinary usage of the English term 'allow'. For the notion of *allowing harm*, as I have defined it and as I shall argue below, is a morally significant notion that has a synergistic relationship with *doing harm*; and this

6. See, e.g., Grice (1975).

remains so even if the notion does not match up completely with the ordinary usage of the English expression. In any case, though, I suggest that the expression ⌜allowing ϕ⌝, as it is ordinarily used in English, *does* fit with my proposed definition — that is, it does not preclude ⌜doing ϕ⌝ or ⌜ϕ-ing⌝, and any temptation to think otherwise is, I suggest, best explained by conversational implicature, and not by the semantics of 'do' and 'allow'.

To define *doing* harm *while allowing it*, we can now simply take our earlier definition of allowing harm and add the condition that the agent *does* the harm of the relevant magnitude that was actually suffered.

> *A* does harm of magnitude *m* to *V* while allowing it $=_{\text{def}}$ (i) *V* suffers a token harm *H* of magnitude *m*, (ii) *A does* this harm to *V*, and (iii) when she did this harm, *A* had an option such that had she taken it, *V* would neither have suffered *H* nor have suffered any harms of total magnitude *m* (or greater) instead of *H* that he didn't actually suffer.

So far, I have argued that, when we do harm of some magnitude to someone, we *usually* also allow them to suffer that magnitude of harm. In §2.2 below, however, I argue that this is not always the case.

2.2. Merely Doing Harm

Let us define *merely doing* harm as doing harm of some magnitude in a way that does not constitute doing harm of that magnitude while allowing it. Is it possible to *merely* do harm? What would *merely* doing harm be like?

Given the definition of doing harm of some magnitude *while* allowing it specified above, to merely do harm is to do some harm while lacking any alternative option such that had we chosen it instead, the harm we actually do (or harm of equal magnitude taking its place) wouldn't have occurred.[7] Thus, we *merely* do harm when we do harm to someone who is, relative to what is in our control, doomed to suffer that harm, or some equivalent harm, no matter what we choose. As I shall argue now, *merely* doing harm is relatively *easy* to justify — easier, in fact, than *merely* allowing harm. And so, perhaps surprisingly, doing harm carries little or no moral weight on its own. Here is an illustrative example of *merely* doing harm:

> *Assured Crushing*: A machine will release a boulder on Tara at time *t*, fatally crushing her, unless a certain button is pressed before *t*. Bob can stop this machine from releasing the boulder only by pressing this button. But pressing this button will trigger a different machine to release the boulder on Tara at *t*, fatally crushing her in just the same manner. So whatever Bob does, the boulder will fall on Tara at *t*, causing her to be fatally crushed in the same manner either way.

If Bob presses the button himself, *he* will have killed Tara. Why? Because *he* will have initiated the sequence that culminates in her death. But Bob does not allow Tara to suffer harm. Why not? First, to simplify, let us say that Tara suffers just one relevant harm — being crushed by the boulder. Bob lacked the ability to do something such that, had he done it, Tara would not have been crushed by the boulder. Thus, there is no relevant harm to Tara such that, had Bob taken a different option, Tara would not have suffered it or an equivalent harm in its place. With respect to Bob's options, in other words, Tara was doomed to suffer the same magnitude of harm from being crushed by the boulder.

Now I don't think that it ultimately matters whether Tara is doomed to suffer the *very same* token harm. What seems to matter is whether Tara is doomed to suffer *either* the harm she actually suffers, *or an equivalent* harm she didn't actually suffer. To motivate that suggestion, first consider how we would morally evaluate Bob's conduct in *Assured Crushing*. Suppose Bob presses the button. In that case Bob does harm

7. If we look to the above definition of *doing harm while allowing it*, we *merely* allow harm when we satisfy all but the second condition, and we *merely* do harm when we satisfy all but the third condition.

to Tara at *t*, but does not allow harm to Tara. Does he act wrongly? The answer is not immediately obvious. It does seem that we would want *some* reason from Bob as to why he pressed the button. That may suggest that Bob needs *some* justification to press the button. On the other hand, though, the considerations that would permit Bob to push the button in *Assured Crushing* could be much weaker than those required if the boulder were *not* otherwise going to be dropped. That is, Bob is morally permitted to push the button in *Assured Crushing* if doing so is necessary to save someone else's life, but the prospect of saving a life would not be enough to justify pushing the button if the boulder was not otherwise going to be dropped on Tara. (That would be to kill one non-doomed innocent in order to save another, which I take to be morally wrong in general.) But in *Assured Crushing*, it strikes me that Bob would be permitted to push the button even to prevent a much lesser harm for someone else (e.g., a broken leg).

Now notice that these judgments seem to hold even if Bob's pressing the button were to cause Tara to suffer a different harm — so long, at least, as we are sure the total magnitude of harm to Tara remains the same. Suppose pressing the button caused Tara to be killed by a trolley at *t* instead of by a boulder, and suppose that being killed by the boulder or the trolley would produce the same total magnitude of harm for Tara (let's imagine that she is unconscious and will feel no difference either way).[8] This change does not seem to make a moral difference; it seems that Bob would still be permitted to press the button in order to save someone else's life, or even to prevent a broken leg, even though in doing so he does lethal harm to Tara while having the option to do something else such that Tara would not have suffered *that* very same harm. Thus, the morally relevant notion of allowing harm seems to be tied to the total magnitude of harm someone suffers: either from the very harms she actually suffers, *or* from equivalently severe harms she didn't actually suffer but would have suffered instead.[9] In the trolley-versus-boulder variation on *Assured Crushing*, Tara is doomed to suffer the same magnitude of harm no matter what Bob does. Bob's pressing the button simply replaces one harm with an equally severe harm, and likewise seems relatively easy to justify.[10]

I submit, then, that even if there is some moral pressure against doing harm without allowing it, this moral pressure is relatively easy to override. In fact, it is weaker than the moral pressure against *merely* allowing harm. For compare *Assured Crushing* with the following case: suppose the machine were set up so that pressing the button would save Tara's life; it would prevent the boulder from killing her *without* releasing any other threats. Bob would need more justification to justify *refraining* from pushing this button than he would to justify pushing the button in *Assured Crushing*. The prospect of preventing a broken leg would *not* be enough to justify Bob's refraining from pressing this button, if it would save Tara's life. For when we must choose between saving a life and preventing a broken leg, I take it that we must save the life. Again, though, it would be permissible in *Assured Crushing* for Bob

8. Her being conscious might introduce unwanted noise into our judgments, as we might have a recalcitrant doubt that being killed one way really is just as harmful as being killed in the other.

9. An anonymous referee asks, what if, by doing some harm to Tara, one thereby prevents Tara from suffering an equivalent harm in the distant future, which Tara was otherwise doomed to suffer? Does it matter whether the harm one replaces happens at or around the same time? For present purposes I shall remain neutral on this issue and explain how to tweak my proposal to accommodate our moral judgments whichever way we settle the question. We can require that, for some harms to count as happening *instead* of *H*, they must happen at or around the same time as *H* would have happened. All examples in this paper are consistent with a narrower, time-sensitive notion of *allowing harm* that incorporates this more stringent requirement; nothing I shall argue for here either rules out or requires this narrower definition of 'allowing harm'.

10. If one were to insist that it *does* matter whether the harm Bob causes Tara by pressing the button is *the very same harm* she would have suffered otherwise — so that pressing the button is easier to justify in the original version of *Assured Crushing* compared the trolley-versus-boulder variant — I could still make the same points about the relationship between doing and allowing harm. I would simply need to contrive the examples in the next section such that the relevantly doomed individuals are doomed to suffer the same exact harm no matter what. I have opted not to do so simply because this move would further complicate examples that are already complicated.

to press the button in order to prevent someone from breaking a leg. So doing harm without allowing it — that is, *merely* doing harm — is actually *easier* to justify than *merely* allowing harm! This claim may seem at odds with the DDA, but it need not be. It is perfectly coherent to accept the DDA, at least in a form suitably clarified, while also accepting that, on its own, doing harm carries *less* moral weight than allowing harm. We shall see how this can be so in §§3–4.

Let us sum up what we have established so far. Most times when an agent does harm to someone, she does harm while allowing it. But it is possible for an agent to do harm without allowing it (though this is uncommon) and also possible for her to allow harm without doing it. Thus, doing and allowing are two separable ways we can relate to harms, and relating to a harm in one way does not preclude our relating to the same harm in the other way. There may be moral pressure against relating in either way to harms. However, the moral pressure against merely doing harm, i.e., doing harm without allowing it, is weaker (i.e., easier to override by other considerations) than the moral pressure against merely allowing harm, i.e., allowing harm without doing it.

3. A Moral Synergy Between Doing and Allowing Harm

It may be tempting to conclude that what really matters in the cases above is only the difference we make for the worse to those who suffer harm. That is, we might think that the moral status of an act is determined by the comparison between how well everyone would fare if the act were performed, and how well everyone would fare had some alternative act been performed instead. Or perhaps, in a less consequentialist vein, we might think that the justification required to permissibly do harm is determined by how poorly off the harm-doing makes the harmed subject, in comparison with how that subject would have fared had the act not been performed. Because in *Assured Crushing* things go just as well for everyone whether Bob pushes the button or not, either view could explain the judgment that Bob is permitted to push the button even though doing so kills Tara.

I do not think that either view is correct, however. Reflection on further cases suggests instead that there is an amplificatory relationship between doing harm and allowing harm — but only when we do harm while allowing it to the same subject.

As a baseline, start with the following case, where one *merely* does harm:

> *Three Trapped (One Doomed)*: A trolley is heading for Tara, Ursula and Victor, who lie unconscious on the tracks. If Bob does nothing, all three will be killed by the trolley. No matter what Bob does, Victor will be crushed to death before he wakes up; Bob cannot wake any of the victims. But Bob can roll a boulder over Victor, crushing him but preventing Tara and Ursula from being crushed by the trolley.

Figure 1: Three Trapped (One Doomed)

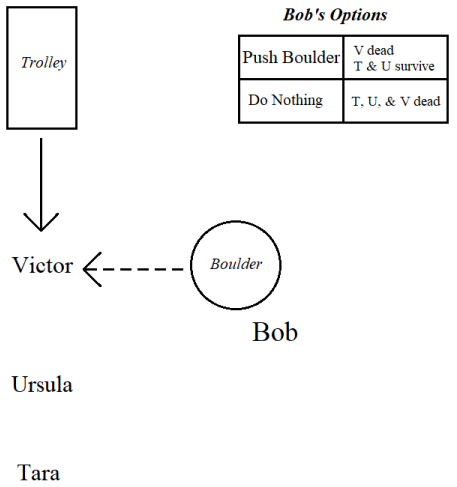

As with the earlier variation on *Assured Crushing*, I think that Bob is permitted to crush Victor in *Three Trapped (One Doomed)*. Victor is doomed to be crushed to death before he wakes up, no matter what Bob does. If Bob does nothing, then Tara, Ursula and Victor will all die. If he rolls the boulder, *he* kills Victor, but at least Tara and Ursula will be saved. Bob may need a reason to push the boulder, but in this case the prospect of saving two lives provides enough reason to justify pushing it.

Now let's contrast *Three Trapped (One Doomed)* with a few other cases. First consider Philippa Foot's *Rescue* II case,[11] which I have modified slightly:

> *Rescue* II: Five people are in imminent, lethal danger. To save them you must drive to them immediately and administer life-saving medicine to each of them. Unfortunately, in order to reach them before they die, you would have to traverse a narrow road where one person is trapped. Thus, you can either do nothing, letting the five die, or you can proceed to rescue the five, running over and killing the one person trapped on the road.

On Foot's view it is impermissible to run over the one person trapped in the road, despite the fact that doing so would enable us to rescue five others. Let's grant that Foot is correct. This verdict stands in striking contrast to *Three Trapped (One Doomed)*. In *Rescue* II it seems wrong to kill one to save five. In *Three Trapped (One Doomed)*, it seems permissible to kill one to save five. What explains the difference? Perhaps, one might think, it is the fact that the one trapped in the road in *Rescue* II would not be under any threat of imminent death if we abstained from saving the five. In *Three Trapped (One Doomed)*, by contrast, Victor will suffer imminent death no matter what we do.

This explanation seems to be on the right track, but by itself I do not believe that it is fully satisfactory. For consider the following variant on *Rescue* II:

> *Rescue* III: Things are as they were in *Rescue* II, except the one person trapped in the road is about to be killed by some independent threat. However, you can very easily stop and free her, thereby saving her life. But doing so will preclude you from saving the five further down the path due to time constraints (in fact, the only way to save the five is to run over and kill the trapped person, just as in *Rescue* II).

In *Rescue* III you have three options: (i) do nothing, and all six people die, (ii) save the one person trapped in the road, letting the five further down the road die, or (iii) run over the person in the road, killing her and saving the five people further down the road. Given that we hold on to the verdict that it's wrong to run over the trapped person in *Rescue* II, it seems clear that option (ii) is the only permissible option in *Rescue* III. It does not matter that the person trapped in the road requires your assistance to survive; it still seems wrong to run her over, just as it seems wrong to run over the trapped person in Foot's *Rescue* II case. What appears to explain the difference between *Three Trapped (One Doomed)* and *Rescue* III is that, in *Rescue* III, you are *able* to do something that would result in survival for the person on the road, while in *Three Trapped (One Doomed)* you are not able to do anything that would result in Victor's survival. This result suggests that doing harm becomes harder to justify not only when abstaining would have resulted in less harm for the victim but also when we have some *option* that would have resulted in less harm for the victim (i.e., when we simultaneously *allow* the victim to suffer harm).

11. Foot (1984: 179).

Given this verdict, let us consider a variant of *Rescue* III, this one closer to *Three Trapped (One Doomed)*:

> *Convergence*: Two trolleys are out of control and are heading toward people who, through no fault of their own, lie temporarily unconscious on the tracks ahead. One trolley is heading for Tara, Ursula and William, and the other toward Victor. If Bob does nothing, all four people will be killed by the trolleys. Bob can roll a boulder to the tracks behind William, thereby preventing Tara and Ursula from being killed by the trolley. But if Bob does this, the boulder will crush Victor and kill him en route to saving Tara and Ursula. Bob can also save Victor from the trolley by shouting loudly, thereby waking him (but not the others) and causing him to move himself off the tracks. If Bob does this, however, he will be unable to move the boulder in time to save Tara and Ursula from the trolley. Tara and Ursula are too far away for Bob to wake them.

Figure 2: Convergence

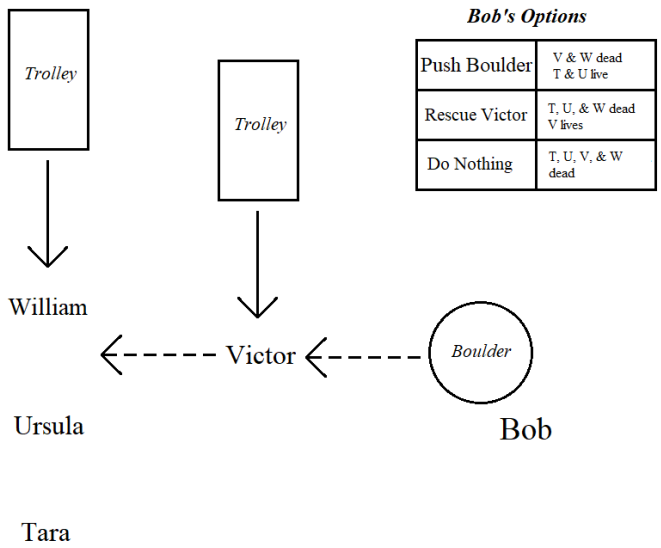

(Note that the reason for including William in the case, for whom Bob can do nothing, will become apparent when we compare *Convergence* with *Divergence* below.)

I suggest that Bob is not permitted to crush Victor with the boulder in *Convergence*, just as we would not be permitted to run over the person trapped in the road in *Rescue* III. Given that Victor is not doomed in this case, as Bob can easily save him by shouting, Bob is not permitted to kill Victor in order to save Tara and Ursula. Harming one to save two from equivalent harms is permissible if, as in *Three Trapped (One Doomed)*, the one is doomed to suffer the same total amount of harm at the same time no matter what one chooses. But killing one to save two is not permissible if the agent can somehow avert the one's death. In that case, by killing one the agent *does* allow the one to suffer harm. Bob is not permitted to kill Victor in *Convergence* because he

could have saved Victor, and thus, in killing him, he would kill Victor *while* allowing him to die. Among the remaining options — easily rescuing Victor, or doing nothing — I take it that the only permissible option is to rescue Victor.

This result has a number of upshots. By pushing the boulder over Victor in *Convergence*, Bob need not make any difference for the worse for Victor; Victor would have suffered the same total amount of harm had Bob done nothing. And even supposing that Bob be so disposed that, had he not pushed the boulder over Victor, he would instead have done nothing — and thus, had Bob not pushed the boulder, Victor would have suffered the same total amount of harm — Bob would still not be permitted to push the boulder. What appears to matter when an agent does harm is not necessarily the difference the agent makes in comparison to what would have happened had the agent done nothing, or even in comparison to what would have happened had the agent not acted as she did. What matters in determining the wrongness of doing harm is whether an alternative was *available* for the harm-doer such that, had the harm-doer taken that alternative instead, the harmed subject would have suffered less total harm.[12]

When we choose between allowing harm to one and allowing harm to another, all else being equal, we are permitted to choose either way. But *Convergence* suggests that when our choice is between doing harm *while* allowing it to one, and merely allowing harm to two others (all else equal), we are forbidden from doing harm *while* allowing it to the one and must instead allow the others to be harmed. This remains so even if by doing harm *while* allowing it to one, the agent leaves the one just as badly off as she would have been had the agent done nothing — or in comparison to how well off she would have been, had the agent not performed the action she did.

But this reasoning leaves open a further question. Does the stringent prohibition against doing harm *while* allowing it to a subject simply stem from the fact that, when we do harm while allowing it, we perform an act that is *pro tanto* wrong in two separate ways, where each way *independently* raises the threshold of justification? Suppose for the sake of simplicity that, if stealing an apple needs to produce X amount of good in order to be justified, and knocking someone down needs to produce Y amount of good in order to be justified, then stealing an apple by pushing someone down needs to produce X + Y amount of good in order to be justified.[13] Is the stringent prohibition against doing *while* allowing harm to a subject simply additive in this way?

Reflection on further cases suggests that it is not.[14] If it were, then we should expect that Bob would also be forbidden from crushing William in the following case.

> *Divergence*: A trolley is heading for Tara, Ursula and William, who lie unconscious on the tracks. Off to the side another trolley is heading for Victor. If Bob does nothing, all four will be killed by trolleys. No matter what Bob does, William will be crushed to death before he wakes up. But Bob can roll a boulder over William, crushing him but preventing Tara and Ursula from being crushed by the trolley. Alternatively, Bob can run to Victor and wake him by shouting loudly, but then Tara, Ursula and William will all be killed by the other trolley.

12. One caveat here: it may matter whether the available alternative in question is morally permissible. For example, if Bob could save Victor only by means of killing an innocent bystander, it may then be permissible for Bob to roll the boulder over Victor after all; perhaps the moral weight of Bob's allowing Victor to die depends on whether Bob had an ability to *permissibly* do something that would have resulted in Victor not dying. If he doesn't, then perhaps, morally speaking, it is as if Bob doesn't count as allowing Victor to die. Thanks to Peter A. Graham for bringing this point to my attention.

13. Cf. Kagan, 1988.

14. I thank Peter A. Graham for his extensive discussion, feedback, and suggestions concerning the development and refinement of the *Divergence* and *Convergence* cases that follow below.

Figure 3: Divergence

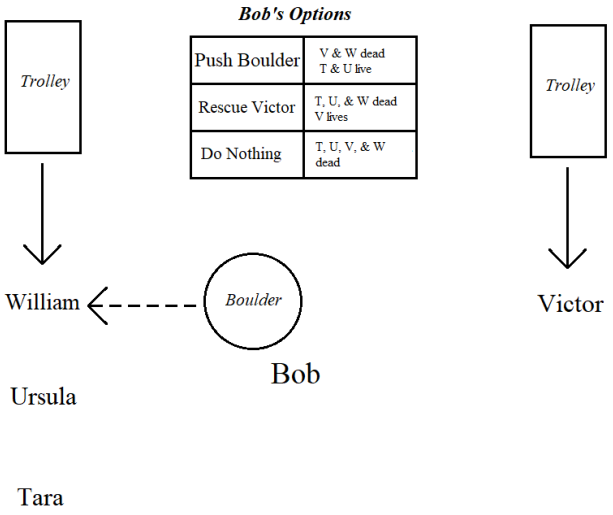

I suggest that it is permissible for Bob to push the boulder onto William, precluding him from saving Victor, in order to save Tara and Ursula. *Divergence* is just like *Three Trapped (One Doomed)*, except that refraining from pushing the boulder would give Bob an option to save one other person off to the side (here, Victor). Since William is doomed to suffer the same magnitude of harm regardless of how Bob acts, it seems that Bob is permitted to push the boulder in order to save a greater total number of lives, just like he is in *Three Trapped (One Doomed)*.[15]

Notice that, just as in *Convergence*, by pushing the boulder in *Divergence* Bob does lethal harm to one and allows lethal harm to another. Consider the overall consequence profile of Bob's options in terms of who lives and who dies, as well as the number of people Bob kills and

15. Note that I do not assume that Bob is morally *required* to crush William; all that matters for present purposes is that he is permitted to do so.

allows to die. In these respects, the two cases are exactly the same. In *Convergence* and *Divergence*, Bob can either (i) do nothing, (ii) rescue Victor or (iii) push the boulder. Doing nothing results in Tara, Ursula, Victor and William all dying (Bob kills no one but allows Victor, Tara and Ursula to die — because he could have saved them — although he does not allow William to die, because he lacks the ability to save William). Rescuing Victor results in Victor's survival but the deaths of everyone else who's trapped (Bob kills no one but allows Tara and Ursula to die, although he again does not allow William to die because he lacks the ability to save William). Pushing the boulder results in Tara and Ursula surviving but Victor and William dying (Bob kills one and allows one to die). What is the difference, then? In *Divergence*, Bob kills a person other than the one he allows to die; he kills William without allowing William to die, and he allows Victor to die without killing Victor. In contrast, in *Convergence* he both kills *Victor* and allows *Victor* to die.

This, I suggest, is why Bob is permitted to push the boulder in *Divergence* but not in *Convergence*. Doing and allowing harm is harder to justify when we do the very harms in virtue of which we allow a magnitude of harm to be suffered. Even though Bob does lethal harm to one person and also allows lethal harm to one person in *Divergence*, the fact that the one Bob kills is different from the one he allows to die seems to make his action more easily justified: that is, justified by the prospect of saving two other people's lives.

There appears, then, to be a synergistic relationship between doing and allowing harm.[16] All else being equal, when we do harm to someone while allowing them to suffer the harm we did to them, that action is harder to justify than when we merely do harm to one and merely allow harm to another. Either doing harm to someone amplifies the justification needed to allow them to suffer harm, or allowing harm to someone amplifies the justification needed to do harm to them. Or both may amplify each other.

16. Or they form a kind of Moorean organic unity. See Moore (1903).

So there is a moral synergy between doing and allowing harm. But we should also wonder *why* this synergy exists. I shall now sketch two possible explanations.

4. Two Possible Explanations

4.1. Amplifying the Moral Weight of Allowing Harm

Let's begin with an explanation for why doing harm could make it harder to justify allowing the total harm to which it contributes.

In general, some justification is needed in order to permissibly allow harm. *More* justification is needed to permissibly allow harms for which one has a special duty not to allow. For example, more justification is needed to allow one's own child to suffer harm than to allow another's child to suffer the same harm. If we must choose between saving our own child from harm and saving another's child from equivalent harm, we must save our own child; we owe it to her to save her. This is perhaps because, as caretakers, we take responsibility for those under our care, which gives us a special obligation not to allow them to suffer harm.[17]

Perhaps doing harm is another way to acquire a special obligation not to allow harm. By creating a threat of harm, an agent thereby becomes morally responsible for what that threat goes on to do — including the harms it causes. Those who launch threats of harm against others thus acquire special obligations not to allow them to suffer harm from the threats they launch, over and above their general duty not to allow others to be harmed. This idea would also account for the intuitive judgment discussed, e.g., in Hanna (2014), that after a threat has been launched, and one faces a choice between saving one's future victim from this threat or saving two unrelated others from equivalent threats, we should save the one from our own threat.[18]

When Bob allows Victor to die in *Divergence*, he allows Victor to suffer harm that he has no special duty to prevent. This is *pro tanto* wrong, because in general there is a *pro tanto* duty not to allow others to suffer harm. But it is ultimately justified by the fact that the only way to save Victor requires Bob to allow both Tara and Ursula to die instead. In *Divergence*, though Bob introduces a threat culminating in William's death — and thereby acquires a special duty not to allow William to suffer harm from this threat — Bob does *not* allow William to suffer harm. This is because he lacks the ability to avert it; he could never have done anything else such that, had he done it, William wouldn't have suffered the harm he actually suffered, or any harm of equal magnitude. Since Bob does not allow William to die, he *a fortiori* does not flout any duty not to allow William to suffer harm.

In *Convergence*, by contrast, Bob is not permitted to allow Victor to die. His option to push the boulder on Victor is morally unacceptable because it *does* flout a duty not to allow Victor him to die. As with William in *Divergence*, Bob's pushing the boulder introduces a threat to Victor's life and thereby gives Bob a special obligation not to allow Victor to suffer total harm from that threat. But unlike *Divergence*, where Bob does not allow William to die, in *Convergence* Bob has the ability to prevent Victor's death. Thus Bob does flout a special duty not to allow Victor to die when he kills him while also allowing him to die. Although Bob could save two others by killing Victor while also allowing him to die, this is not enough to justify Bob's allowing Victor to die. By launching the lethal threat toward Victor, Bob thereby acquired a special obligation not to allow Victor to die, but he lacks such a special obligation to Tara and Ursula. If the only way Bob can satisfy this special obligation to Victor is not to roll the boulder toward Victor in the first place, then he is not permitted to do so.

To repeat, on this explanation, *allowing* harm does the direct moral work, and *doing* harm is an aggravating factor; it alters the moral status of allowing harm, in particular by increasing the threshold for its

17. Of course, if those in our care suffer harm in a way we are utterly powerless to stop, we have not wronged them, because we have not allowed them to suffer harm. If we never had the ability to stop something from happening, it is false that we allowed it to happen.

18. Cf. Unruh (2021) for a similar proposal on which doing harm. or rather launching a threat of harm, amplifies the moral pressure not to allow harm by giving rise to a special duty not to allow that harm.

justification. Note, however, that for this story to work on its own in explaining the moral synergy between doing and allowing harm, we may need to add that different kinds of special duties to prevent harm have different degrees of moral weight. For we might hold that we are not permitted to kill an innocent bystander to save our own child's life, yet we want to say that we have special duties to prevent harm to children under our care.[19] If the wrongness of killing the bystander is explained entirely by the fact that, in doing so, we allow a harm for which we have a special duty not to allow, then why wouldn't killing the bystander while allowing her to die be morally *equivalent* to allowing our child to die? In either case one allows a death that one had a special duty to prevent. If the strength of special duties to prevent harm comes in degrees, however, we can then explain why the duty not to kill the bystander while allowing her to die wins out: we can do so by holding that doing harm gives us a *stronger* special duty not to allow said harm, compared with the special duties caregivers have not to allow harm to come to those under their care.

Of course we would also want to know why this is so, and for now, I leave it open how the story would go and whether it could be justified on independent grounds. In the following section, though, I offer an alternative that could either supplement or replace the story just considered.

4.2. Amplifying the Moral Weight of Doing Harm

In this section I explore an alternative way of explaining the moral synergy between doing and allowing harm, on which allowing a harm that we also do makes it harder to justify doing that harm, in comparison with *merely* doing the harm.

Recall *Assured Crushing*, discussed in §2.2. There I suggested that Bob might need *some* justification to permissibly push the button. The story discussed in §4.1 could not explain why that would be so. If doing harm were wrong only insofar as it amplifies the wrongness of allowing harm, then it would not be wrong at all for Bob to harm Tara by pushing the button in *Assured Crushing*. This is because Bob does not allow any harm by doing so. Perhaps, then, the wrongness of doing harm while allowing it may not be fully explained by the fact that, in general, doing harm modulates the stringency of our duty not to allow harm. There may be a distinctive wrongness associated with doing harm irrespective of the wrongness of allowing harm.

So must Bob have some justification to push the button in *Assured Crushing*? Let's suppose that, in general, we have the right not to have harm done to us. If rights-infringements are always *pro tanto* wrong, then Bob would need *some* justification to push the button, because doing so would infringe Tara's right not to have harm done to her and thereby be *pro tanto* wrong to some (small) degree.

Arguably, though, such rights-infringements can have varying degrees of moral weight. For example, using a pencil to painfully poke another's hand without her consent would infringe her right not to have harm done to her. It would seem to be *pro tanto* wrong. Using a sharp knife to stab someone's hand without her consent would also infringe that right, but it is *pro tanto* wrong to a greater degree. Why? The explanation may lie in a general feature of rights-infringements captured by Judith Jarvis Thomson's Aggravation Principle,[20] which we can state roughly this way: the worse it is for a victim that a right of hers is infringed, the more gravely the victim is wronged by that rights-infringement. Or rather, I wish to consider a narrower formulation: the more harmful it is for a victim that a right of hers is infringed, the more gravely the victim is wronged by that rights-infringement.[21] When some harm that someone suffers is not one we allowed to hap-

19. In a similar vein, notice that if someone wrongfully pushes an innocent in front of a trolley and subsequently must choose between saving that innocent or their own child, we might think that they must save the innocent they put in harm's way. Thanks to Peter A. Graham for suggesting this case.

20. Thomson (1990: 267).

21. We might also think that rights-infringements can become more serious by closing off greater benefits to the victims, although perhaps to a lesser degree than by being harmful to them. I leave this issue open. I thank an anonymous referee for pressing me to be more explicitly sensitive to the distinction between harming and failing to benefit in discussing the Aggravation Principle.

pen, the one who suffers that harm would have suffered the same total magnitude of harm no matter what we chose. Thus, when we do harm without allowing it, our doing the harm does not result in her suffering more harm than she would have, had we acted differently. Cases like *Convergence*, I think, suggest something interesting about the Aggravation Principle: it ought to be spelled out not in terms of how much harm a victim would have suffered had the aggressor done nothing, nor in terms of how much harm would have been suffered had the aggressor not infringed the right — but rather in terms of how much harm arising from the rights-infringement the victim *could* have been spared, had the aggressor done something else she was able to do.[22]

This explanation, then, combines two ideas: (i) that harm-doing infringes some right which harm-allowing would not in itself infringe, and (ii) that, in general, rights-infringements carry more moral weight the more harmful it is for the victim that the right was infringed. If we infringe someone's right without allowing them to suffer harm, the victim could not have been spared any magnitude of harm no matter what we did. If we infringe a right by doing harm *while* allowing it, however, we act in such a way that the victim could have been spared some harm arising from the rights-infringement had we acted differently — meaning that it is harmful for the victim that we chose to infringe the right rather than doing something else we could have done. Thus, if we grant both the Aggravation Principle and that doing harm infringes a right, we can readily explain how allowing harm to someone amplifies the moral weight of doing the harm to her that contributes to this overall harm.

In *Assured Crushing* let's suppose Tara has right against Bob that he not push the button (given that pushing the button would do harm to her). For that reason, it may be *pro tanto* wrong to some (perhaps small) degree for Bob to press the button. But if, unlike in *Assured Crushing*, Bob could prevent Tara from being killed, then pushing the button would make Tara suffer more harm in comparison to what would have happened had Bob taken a different option (i.e., saving Tara's life). In that case, the Aggravation Principle would kick in. It would increase the gravity of the infringement of Tara's right, and thereby increase the threshold needed to justify pressing the button.

Now how does this story account for the difference between *Convergence* and *Divergence*? Bob does infringe a right not to be harmed by pushing the boulder in both *Divergence* and *Convergence*. But this right-infringement takes more to justify in the latter case than in the former, because it leads to more harm for the victim (Victor) in *Convergence*, compared with how much he could have suffered had Bob acted in some other way he was able to act; but it does not lead to more harm for the victim (William) in *Divergence* in comparison with how much he could have suffered, had Bob acted in some other way in which he was able to act. Given that, in *Divergence*, Bob does not infringe any right of *Victor*'s no matter what he does, there is no rights-infringement in *Divergence* that is nearly as grave as Bob's infringement of Victor's right when in *Convergence* he pushes the boulder onto him. So, in *Convergence* Bob commits a grave infringement of a right by pushing the boulder onto Victor, whereas in *Divergence* he commits a minor infringement of the same kind of right by pushing the boulder onto William while, at the same time, also failing to save one person's life. Saving two people's lives is, I take it, a weighty enough consideration to morally justify such a minor rights-infringement plus a failure to save one other person's life,[23] but it is not a weighty enough consideration to morally justify the grave right-infringement involved in lethally crushing a non-doomed person, as Bob does to Victor in *Convergence*.

22. Cf. Norcross (2005: 165–166), in particular his Bobby Knight case.

23. Compare with a case where Bob must choose between (i) doing nothing, (ii) saving one person's life and (iii) saving two people's lives in a way involving pushing an unconsenting bystander out of his way. Although pushing the bystander in option (iii) infringes a right of the bystander's, this infringement carries a relatively small degree of moral weight because it does not harm her significantly. I take it that Bob would be permitted to take option (iii) in such a case, even though it involves an infringement of a right. Just like in *Divergence*, the rights-infringement is not grave enough to outweigh the moral pressure exerted by the prospect of saving two lives.

5. Conclusion

We can relate to harms in multiple ways. Doing and allowing are at least two such ways, and I have argued that relating in one way to harm neither requires nor precludes relating to it in the other way. Far from being an unusual type of case, *most* cases of harm-doing are also cases in which the harm-doer allows the harm they inflict, thereby allowing some total magnitude of harm to be suffered. Recognizing this fact serves as a steppingstone to help us see the moral synergy between doing and allowing harm: either the moral weight of allowing some total amount of harm is amplified when one also does the harms contributing to the magnitude of harm allowed to be suffered, or the moral weight of doing harm is amplified when one also allows the same subject to suffer harm at the same time — or both.

I then offered two preliminary explanations for this moral synergy. On the first, doing harm amplifies the moral weight of allowing that same harm because it is, in general, harder to justify allowing harm when we have a special duty not to allow it, and doing a harm is a way of acquiring such a special duty. On the second explanation, doing harm to someone is directly *pro tanto* wrong because it infringes a right of hers. Given that rights-infringements are in general graver the more harmful they are for their victims, an infringement of this kind will be less grave when the victim would have suffered just as much harm regardless of what the aggressor did. Thus, harm-doing will carry less moral weight when the harm-doer doesn't thereby allow harm. On this story, allowing harm amplifies the moral weight of doing the harm in question.

These two stories are entirely compatible with each other. I leave it open whether we should accept one or both of them. There is also much more to say about the stories themselves. I leave the details to future work. For now I conclude by suggesting we reframe the DDA to say that it is doing harm *while also allowing it* that is harder to justify than merely allowing harm. We should recognize that it is the convergence of two ways of relating to harm — *both doing* and *allowing* — that makes ordinary harm-doing so hard to justify.[24]

References

Bennett, Jonathan (1995). *The Act Itself*. New York: Oxford University Press.

Foot, Philippa (1967). The problem of abortion and the doctrine of double effect. *Oxford Review*, 5: 5–15.

—— (1984). Killing and letting die. In Jay Garfield and Patricia Hennessey (eds.), *Abortion: Moral and Legal Perspectives*. Amherst: University of Amherst Press, pp. 177–185.

—— (1985). Morality, action and outcome. In Ted Honderich (ed.), *Morality and Objectivity*. London: Routledge and Kegan Paul, pp. 23–38.

Grice, Herbert Paul (1975). Logic and conversation. In Donald Davidson and Gilbert Harman (eds.), *The Logic of Grammar*. Encino, CA: Dickenson, pp. 64–75.

Hanna, Jason (2014). Doing, allowing, and the moral relevance of the past. *Journal of Moral Philosophy*, 12: 677–698.

Kagan, Shelly (1988). The additive fallacy. *Ethics*, 99(1): 5–31.

Liu, Xiaofei (2012). A robust defense of the doctrine of doing and allowing. *Utilitas*, 24(1): 63–81.

Moore, George Edward (1903). *Principia Ethica*. Cambridge: Cambridge University Press.

Norcross, Alastair (2005). Harming in context. *Philosophical Studies: An International Journal for Philosophy in the Analytic Tradition*, 123(1/2): 149–173.

24. I am deeply indebted to Peter A. Graham for working closely with me to improve this paper over the whole span of several years that I worked on it, providing critical feedback, encouragement, and suggestions at every stage in the development of the paper. I am also deeply indebted to Alejandro Pérez Carballo, Cruz Davis, Maya Eddon, Firdaus Gupte, Justis Koon, Christopher Meacham, Justin Mooney, Jean Pierre Cordero Rojas, Sam Schechter, and some especially helpful anonymous referees; your feedback, encouragement, and discussions were all crucial to the development of this paper, and I am extremely grateful for it. Thank you all very much.

Quinn, Warren (1989). Actions, intentions, and consequences: The doctrine of doing and allowing. *Philosophical Review,* 98(3): 287–312.

Rachels, James (1975). Active and passive euthanasia. *The New England Journal of Medicine*, 292(9): 78–80.

Thomson, Judith Jarvis (1976). Killing, letting die, and the trolley problem. *The Monist*, 59(2): 204–217.

—— (1990). *The Realm of Rights*. Cambridge, MA: Harvard University Press.

Unruh, Charlotte Franziska (2021). Letting climate change. *Journal of the American Philosophical Association*, 7(3): 368–386.

Woollard, Fiona (2015). *Doing and Allowing Harm*. Oxford: Oxford University Press.